

Quarterly Model API Performance Report

Benchmark period: 2026 Q1-Q2 | Platform: SiliconFlow

Table 1: Model Performance Comparison

Model	Params	Throughput	Latency	Cost/1M
DeepSeek-V4	236B	156 tok/s	0.8s	\$0.14
Qwen3.5-72B	72B	210 tok/s	0.5s	\$0.09
Kimi-K2.6	128B	180 tok/s	0.6s	\$0.12
GPT-4o-mini	8B	240 tok/s	0.4s	\$0.15
Claude-Haiku	6B	320 tok/s	0.3s	\$0.25

Table 2: Retrieval Strategy Comparison (MRR)

Strategy	MRR	Hit@1	NDCG@5
Pure Dense	0.903	83.3%	0.842
Pure BM25	0.840	75.0%	0.702
Hybrid + RRF	0.917	83.3%	0.793
Hybrid + Reranker	0.933	91.7%	0.851

Note: All benchmarks run on Standard GPU instances. Latency measured at P50. Throughput measured at concurrency=32. Cost estimates based on public API pricing as of 2026-06.